

Linear Regression

Jim Harner

1/6/2020

sparklyr requires a dplyr compatible back-end to Spark.

```
library(dplyr, warn.conflicts = FALSE)

# load the sparklyr package
library(sparklyr)
# start the sparklyr session
master <- "local"
# master <- "spark://master:7077"
sc <- spark_connect(master)
```

6.3 Concrete Slump Test Regression

Load slump.csv into Spark with spark_read_csv from the local filesystem.

```
slump_sdf <- spark_read_csv(sc, "slump_sdf",
  path = "file:///home/rstudio/rspark-tutorial/data/slump.csv")
head(slump_sdf)
```

```
## # Source: spark<?> [?? x 10]
##   cement slag fly_ash water    sp coarse_aggr fine_aggr slump  flow
##   <dbl> <dbl>   <dbl> <dbl> <dbl>      <dbl>      <dbl> <dbl> <dbl>
## 1    273    82    105   210     9         904        680    23    62
## 2    163   149    191   180    12         843        746     0    20
## 3    162   148    191   179    16         840        743     1    20
## 4    162   148    190   179    19         838        741     3   21.5
## 5    154   112    144   220    10         923        658    20    64
## 6    147    89    115   202     9         860        829    23    55
## # ... with 1 more variable: compressive_strength <dbl>
```

First we need to split slump_sdf into a training and a test Spark DataFrame.

```
slump_partition <- tbl(sc, "slump_sdf") %>%
  sdf_partition(training = 0.7, test = 0.3, seed = 2)
slump_train_sdf <- slump_partition$training
slump_test_sdf <- slump_partition$test
```

The full model is now run.

```
slump_lr_full_fit <- slump_partition$training %>%
  ml_linear_regression(compressive_strength ~ cement + slag + fly_ash + water
    + sp + coarse_aggr + fine_aggr)
summary(slump_lr_full_fit)
```

```
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -5.6280 -1.6192 -0.3183  0.9372  7.1920
##
## Coefficients:
```

```
## (Intercept)      cement      slag      fly_ash      water      sp
## 219.36232986    0.03777496  -0.06065688   0.02819246  -0.31892157  -0.12983604
## coarse_aggr    fine_aggr
## -0.08744781   -0.06805072
##
## R-Squared: 0.8987
## Root Mean Squared Error: 2.507
```

Notice that the model summary does not provide much useful information. We can p-values by by getting a tidy summary.

```
tidy(slump_lr_full_fit)
```

```
## # A tibble: 8 x 5
##   term      estimate std.error statistic p.value
##   <chr>      <dbl>    <dbl>    <dbl>   <dbl>
## 1 (Intercept) 219.      92.4      2.37  0.0207
## 2 cement      0.0378    0.0290     1.30  0.198
## 3 slag      -0.0607    0.0409    -1.48  0.143
## 4 fly_ash     0.0282    0.0301     0.938 0.352
## 5 water     -0.319     0.0927    -3.44  0.00104
## 6 sp        -0.130     0.183     -0.708 0.481
## 7 coarse_aggr -0.0874    0.0355    -2.46  0.0166
## 8 fine_aggr  -0.0681    0.0379    -1.79  0.0778
```

Performance metrics for regression are generally obtained first by getting predictions and then using an evaluator to get a specific metric.

```
slump_lr_full_predict <- ml_predict(slump_lr_full_fit)
slump_lr_full_predict
```

```
## # Source: spark<?> [?? x 11]
##   cement slag fly_ash water sp coarse_aggr fine_aggr slump flow
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 137 167 214 226 6 708 757 27.5 70
## 2 140 1.4 198. 175. 4.4 1050. 780. 16.2 31
## 3 140 128 164 183 12 871 775 23.8 53
## 4 140 128 164 237 6 869 656 24 65
## 5 140. 4.2 216. 194. 4.7 1050. 710. 24.5 57
## 6 140. 11.8 226. 208. 4.9 1021. 684. 21 64
## 7 140. 44.8 235. 171. 5.5 1048. 704 23.5 52.5
## 8 140. 61.1 239. 182. 5.7 1018. 681. 24.5 60
## 9 142 130 167 174 11 883 785 0 20
## 10 143 131 168 217 6 891 672 25 69
## # ... with more rows, and 2 more variables: compressive_strength <dbl>,
## # prediction <dbl>
```

```
ml_regression_evaluator(slump_lr_full_predict, label_col = "compressive_strength",
  prediction_col = "prediction", metric_name = "rmse")
```

```
## [1] 2.50723
```

This would be awkward if want to evaluate a series of models for several metrics.

The model for the lasso with varying values of the regularization parameter λ .

```
slump_perf_metrics <- function(l) {
  slump_train_sdf %>%
```

```

ml_linear_regression(compressive_strength ~ cement + slag + fly_ash +
                     water + sp + coarse_aggr + fine_aggr,
                     elastic_net_param = 1, reg_param = 1)
}

```

First, we Initialize the performance data frames for $\lambda = 0$. Notice that we can get the performance metrics as the components of summary list, which in turn is an element of the fitted list.

```

regParm <- c(0.02, 0.04, 0.06, 0.08, 0.1, 0.12, 0.14)
slump_lr_errors <- data.frame(lambda = 0,
                              r2 = slump_lr_full_fit$summary$r2,
                              rmse = slump_lr_full_fit$summary$root_mean_squared_error,
                              mae = slump_lr_full_fit$summary$mean_absolute_error)
slump_lr_coef <- as.data.frame(slump_lr_full_fit$coefficients)

```

We now calculate r2, rmse, and mae for each of the models.

```

for(l in regParm) {
  slump_lr_fit <- slump_perf_metrics(l)
  slump_lr_errors <-
    data.frame(lambda = l,
              r2 = slump_lr_fit$summary$r2,
              rmse = slump_lr_fit$summary$root_mean_squared_error,
              mae = slump_lr_fit$summary$mean_absolute_error) %>%
    rbind(slump_lr_errors, .)
  slump_lr_coef <-
    as.data.frame(slump_lr_fit$coefficients) %>%
    cbind(slump_lr_coef, .)
}
slump_lr_errors

```

```

##   lambda      r2      rmse      mae
## 1  0.00 0.8987442 2.507230 1.896407
## 2  0.02 0.8970112 2.528594 1.882784
## 3  0.04 0.8948250 2.555291 1.905668
## 4  0.06 0.8944239 2.560160 1.912491
## 5  0.08 0.8938638 2.566941 1.919315
## 6  0.10 0.8931438 2.575633 1.927251
## 7  0.12 0.8922638 2.586218 1.941522
## 8  0.14 0.8912237 2.598671 1.958873

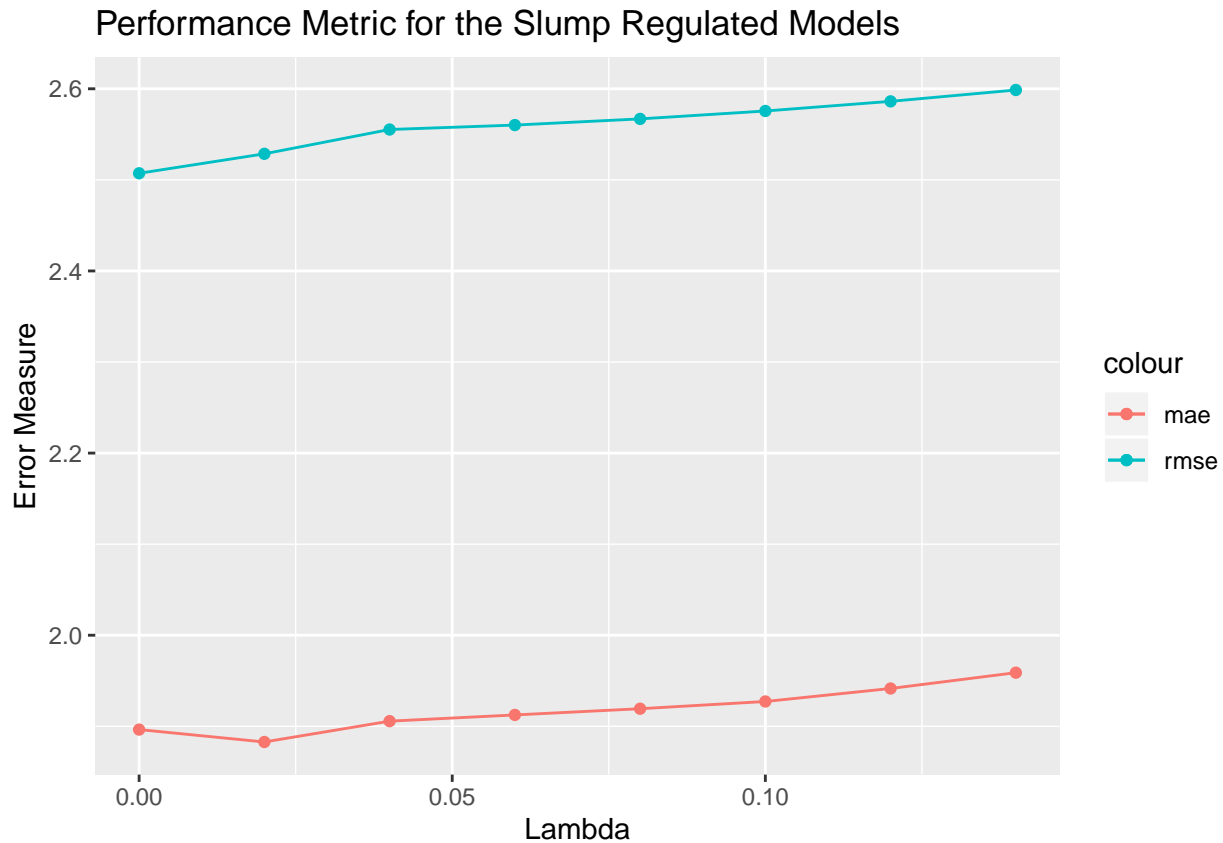
```

Finally, we plot the performance measures.

```

library(ggplot2)
slump_lr_errors %>%
  ggplot(aes(x = lambda)) +
  geom_point(aes(y = rmse, color = 'rmse')) +
  geom_line(aes(y = rmse, color = 'rmse')) +
  geom_point(aes(y = mae, color = 'mae')) +
  geom_line(aes(y = mae, color = 'mae')) +
  ggtitle("Performance Metric for the Slump Regulated Models") +
  xlab("Lambda") + ylab("Error Measure")

```



Based on the performance metrics, it is clear we want `lambda` to be small. However, we also want parsimony.


We now get the parameter estimates as `lambda` increases.

```
names(slump_lr_coef) <- as.character(rbind(c(0.0, regParm)))
slump_lr_coef <- t(slump_lr_coef)
slump_lr_coef
```

```
##      (Intercept)      cement      slag      fly_ash      water      sp
## 0      219.36233  0.03777496 -0.06065688  0.02819246 -0.3189216 -0.129836041
## 0.02   125.22235  0.06625311 -0.01963970  0.05789226 -0.2252704  0.000000000
## 0.04    80.23832  0.07943291  0.00000000  0.07174224 -0.1803432  0.036599089
## 0.06    78.41981  0.07912863  0.00000000  0.07149383 -0.1772967  0.031142392
## 0.08    76.60486  0.07882439  0.00000000  0.07125072 -0.1742739  0.025677415
## 0.1     74.78988  0.07852014  0.00000000  0.07100761 -0.1712511  0.020212383
## 0.12    72.97495  0.07821590  0.00000000  0.07076451 -0.1682284  0.014747354
## 0.14    71.15997  0.07791167  0.00000000  0.07052140 -0.1652056  0.009282288
##      coarse_aggr      fine_aggr
## 0      -0.08744781 -0.068050721
## 0.02   -0.05126774 -0.029635158
## 0.04   -0.03384715 -0.010993751
## 0.06   -0.03296010 -0.010196193
## 0.08   -0.03207201 -0.009399424
## 0.1    -0.03118391 -0.008602646
## 0.12   -0.03029584 -0.007805887
## 0.14   -0.02940773 -0.007009110
```

The lasso trace of the coefficient estimates provides a way of picking the strength of regulation.

```
library(ggplot2)
as.data.frame(cbind(lambda = c(0.0, regParm), slump_lr_coef)) %>%
  ggplot(aes(x = lambda)) +
  geom_line(aes(y = cement, color = 'cement')) +
  geom_line(aes(y = slag, color = 'slag')) +
  geom_line(aes(y = fly_ash, color = 'fly_ash')) +
  geom_line(aes(y = water, color = 'water')) +
  geom_line(aes(y = sp, color = 'sp')) +
  geom_line(aes(y = coarse_aggr, color = 'coarse_aggr')) +
  geom_line(aes(y = fine_aggr, color = 'fine_aggr')) +
  ggtitle("Parameter Trace for the Slump Regulated Models") +
  xlab("Lambda") + ylab("Coeff. Estimate")
```



Over the range of λ , we have 3 features (**cement**, **fly_ash**, and **water**) with consistently non-zero coefficient estimates. Arguably, **coarse_aggr** also deviates from 0. These agree with the model we found by *ad hoc* variable selection in Section 6.1.

At this point we could pick several models to run on the test Spark DataFrame for final selection.

```
spark_disconnect(sc)
```